

Day 2: Linear Regression

Summer STEM: Machine Learning

Department of Electrical and Computer Engineering
NYU Tandon School of Engineering
Brooklyn, New York

Outline

1 Leftovers from Day 1

2 Introduction to Machine Learning

3 Statistics Basics

4 Linear Regression

Exercises: Matrix Multiplication

$$(AB)_{ij} = \sum_{k=1}^N A_{ik} B_{kj} \quad (A^T)_{ij} = A_{ji}$$

■ $X = \begin{bmatrix} 1 & 2 & -1 \\ 1 & 0 & 1 \end{bmatrix}$ $Y = \begin{bmatrix} 3 & 1 \\ 0 & -1 \\ -2 & 3 \end{bmatrix}$ $Z = \begin{bmatrix} 1 \\ 4 \\ 6 \end{bmatrix}$

■ Calculate \underline{XY} , \underline{YX} , $\underline{Z^T Y}$

$$Z^T = [1 \ 4 \ 6]$$

$$Z^T Y = [1 \ 4 \ 6] \begin{bmatrix} 3 & 1 \\ 0 & -1 \\ -2 & 3 \end{bmatrix} = [3 + (-12) \ 1 - 4 + 18] = [-9 \ 15]$$

$\begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix}^{-1} \begin{bmatrix} 5 \\ 13 \end{bmatrix}$

Matrix Inverse

A m by n matrix

$$A^{-1}A = AA^{-1}$$

m has to be equal to n

- Analogy: Reciprocal of a number $\frac{1}{a}a = 1$
- Matrix inverse only defined for square matrix (# rows = # cols)
- $(A^{-1}A = AA^{-1} = I)$. I is called the identity matrix, whose diagonal elements are 1 and other elements are 0.
- Hard to compute by hand, but for 2 by 2 matrix, it is identity

$$[a]^{-1} = \left[\frac{1}{a} \right]$$

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix}^{-1} = \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 \end{bmatrix}$$

Ex: 2 by 2

$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

- Even for a square matrix, the matrix inverse does not always exist. Can you tell when that is the case for 2 by 2 matrices based on the formula given above?

$$ad - bc = 0$$

Matrix Inverse

When is matrix inverse useful? We can use it to solve systems of linear equations!

- Consider the following equations

$$\begin{cases} x + 2y = 5 \\ 3x + 5y = 13 \end{cases} \quad \text{Ex: } \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} = \begin{bmatrix} u_1 + 0 \\ 0 + u_2 \end{bmatrix} = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}$$

- In matrix-vector form

$$\begin{aligned} & \begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix}^{-1} \quad \begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 5 \\ 13 \end{bmatrix} \\ &= \frac{1}{5 \times 1 - 2 \times 3} \begin{bmatrix} 5 & -2 \\ -3 & 1 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix}^{-1} \begin{bmatrix} 5 \\ 13 \end{bmatrix} \\ &= \begin{bmatrix} -5 & 2 \\ 3 & -1 \end{bmatrix} \underbrace{\begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix}}_{I} \begin{bmatrix} x \\ y \end{bmatrix} = \underbrace{\begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix}^{-1}}_{\{}} \begin{bmatrix} 5 \\ 13 \end{bmatrix} \end{aligned}$$

Demo and Exercises: NumPy

Open `demo_vectors_matrices.ipynb`

- Your task: use NumPy functions to compute the exercises we did earlier this morning. Compare the results.

Demo: Plotting Functions

- Generate and plot the following functions in Python:
 - Scatter plot of points: $(0,1), (2,3), (5,2), (4,1)$
 - Straight Line: $y = mx + b$
 - Sine-wave $y = \sin(x)$
 - Polynomial e.g. $y = x^3 + 2$
 - Exponential e.g. $y = e^{-2x}$
 - Choose a function of your own
- Use Wikipedia and Numpy Documentation to search for mathematical formulas and python functions

Looking at our ice-breaker data in spreadsheets

- Columns have labels in the first row
- Collected data (numbers, words) follow below
- Let's export it to a Comma-Separated Values (CSV) file and open it

Outline

1 Leftovers from Day 1

2 Introduction to Machine Learning

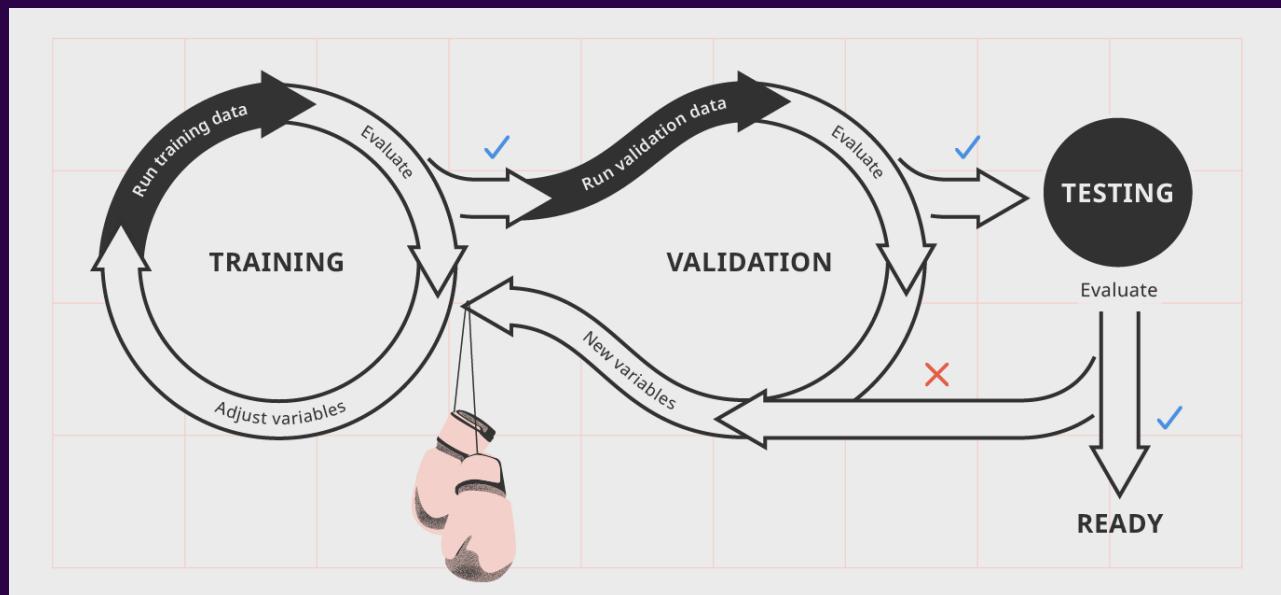
3 Statistics Basics

4 Linear Regression

What is Machine Learning

- Recognize patterns from data
- Make predictions based on the learnt patterns
- A very effective tool where human expertise is not available

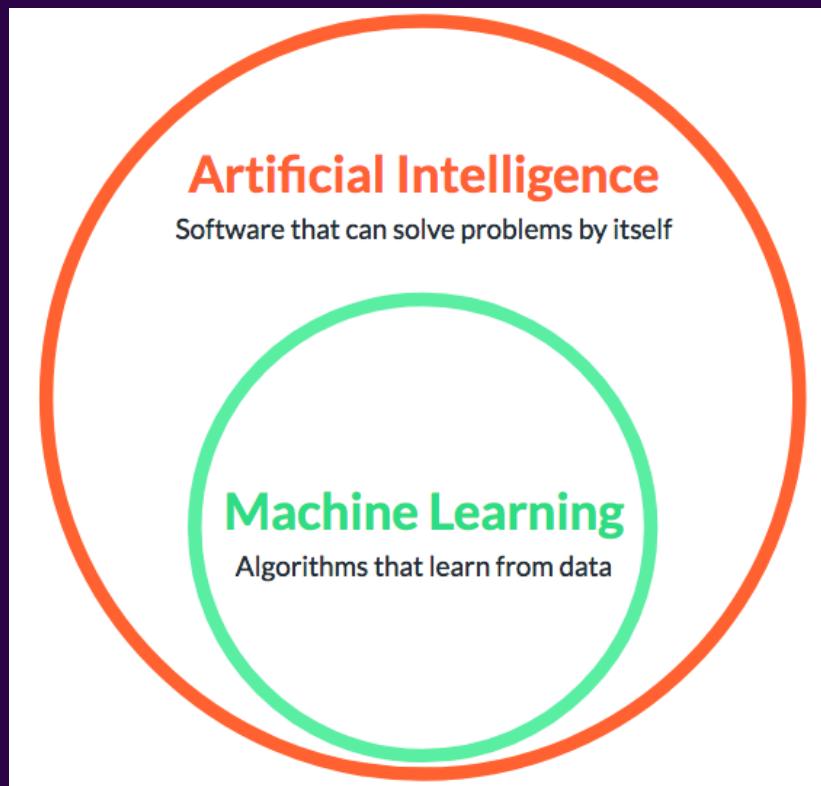
Machine Learning Pipeline



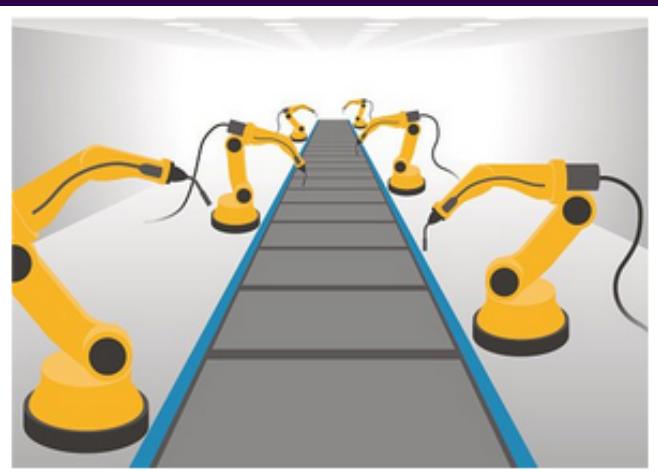
Artificial Intelligence

- Search
- Reasoning and Problem Solving
- Knowledge Representation
- Planning
- Learning
- Perception
- Natural Language Processing
- Motion and Manipulation
- Social and General Intelligence

Machine Learning



Autonomous vs. Automated



Autonomous Example: Self-driving car



■ Waymo Video

<https://www.tesmanian.com/blogs/tesmanian-blog/tesla-autopilot-full-self-driving-fsd-improvements-video>

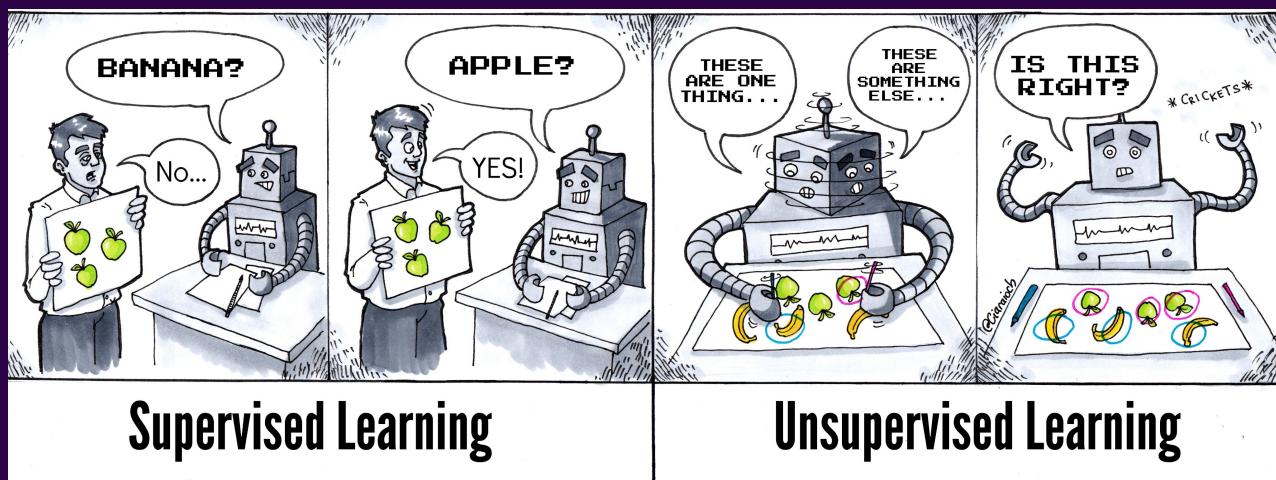
Why is Machine Learning so Prevalent?

- Database mining
- Medical records
- Computational biology
- Engineering
- Recommendation systems
- Understanding the human brain

Why Now?

- Big Data
 - Massive storage. Large data centers
 - Massive connectivity
 - Sources of data from internet and elsewhere
- Computational advances
 - Distributed machines, clusters
 - GPUs and hardware

Supervised Vs. Unsupervised Learning

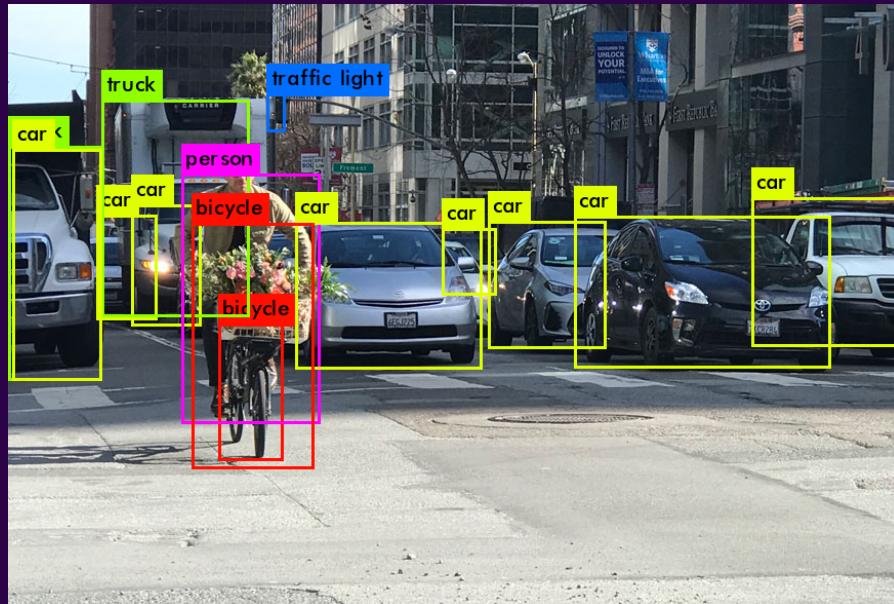


twitter.com/athena_schools/status/1063013435779223553/photo/1 

Supervised Vs. Unsupervised Learning

- The main difference between supervised and unsupervised learning is the existence of a supervisor, which in many cases is in the form of a data label.
- The label of the data is what we want the machine learning algorithm to predict.

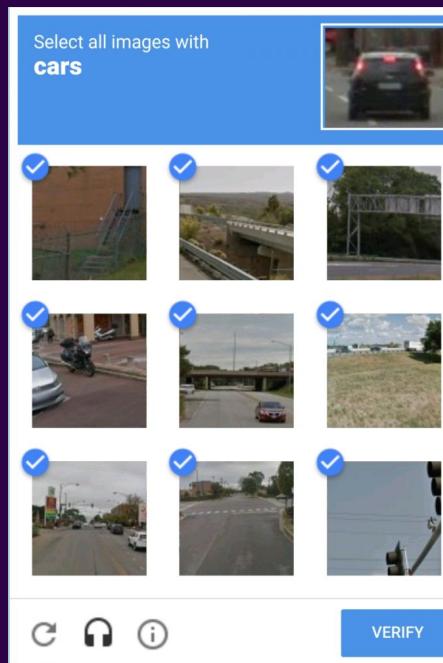
Labelled Data



■ YOLO v2

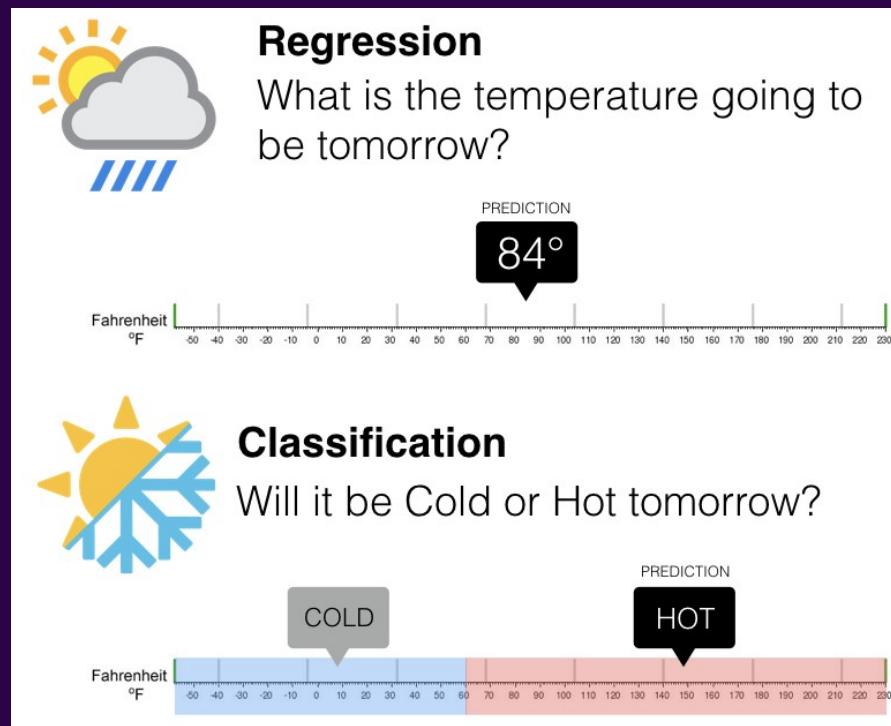
<https://towardsdatascience.com/yolo-you-only-look-once-17f9280a47> NYU TANDON SCHOOL OF ENGINEERING

How labels are generated



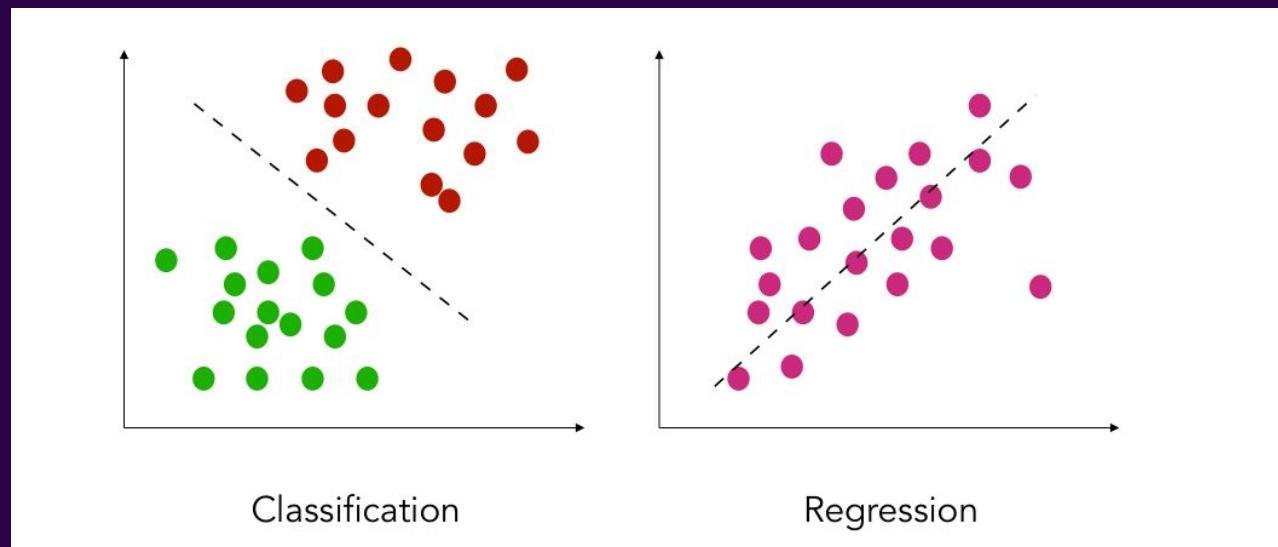
<https://devrant.com/rants/1758134/select-all-images-with-cars-i-did-and-its-not-correct-why-not>

Classification Vs. Regression

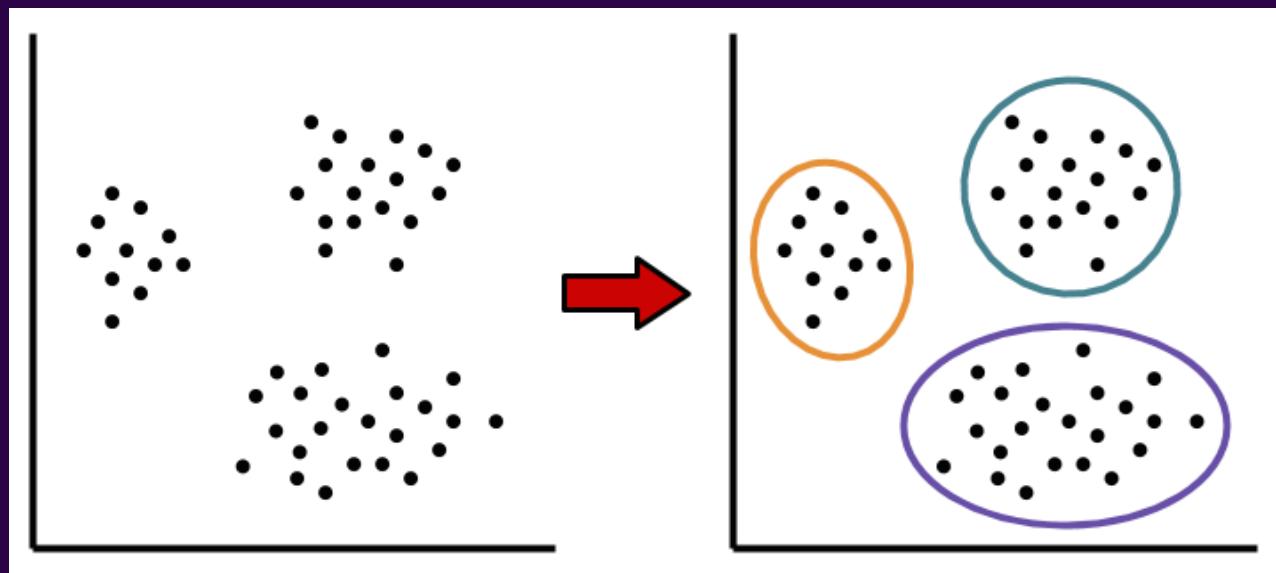


<https://www.pinterest.com/pin/672232681855858622/?lp=true>

Classification Vs. Regression



Unsupervised Learning



source: the dish on science

Outline

1 Leftovers from Day 1

2 Introduction to Machine Learning

3 Statistics Basics

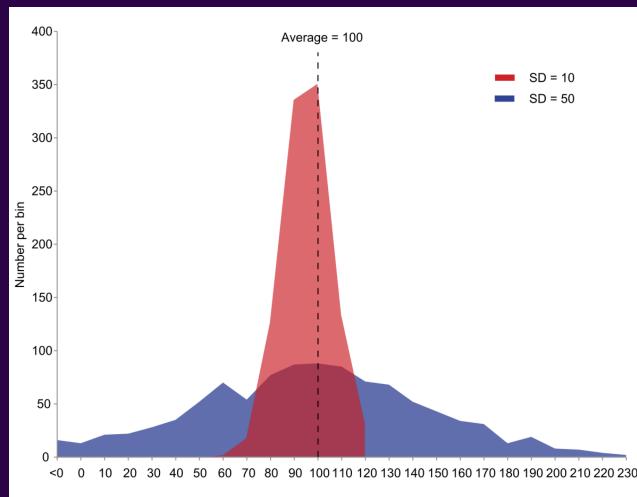
4 Linear Regression

Basic Concepts

- **Mean** (average value): $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$
- **Variance** describes the spread of the data with respect to the mean.
- **Covariance** describes the relationship between two variables.

Variance

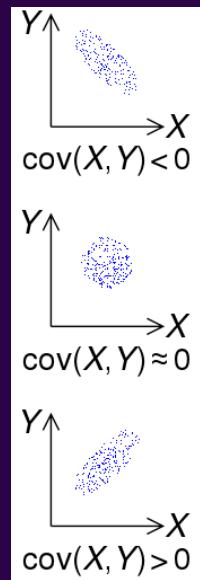
■ Variance: $\sigma_x^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2$



<https://en.wikipedia.org/wiki/Variance>

Covariance

- Covariance: $\sigma_{xy} = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$



<https://en.wikipedia.org/wiki/Covariance>

Outline

1 Leftovers from Day 1

2 Introduction to Machine Learning

3 Statistics Basics

4 Linear Regression

Supervised learning in a nutshell

Given the dataset (x_i, y_i) for $i = 1, 2, \dots, N$, find a function $f(x)$ (model) so that it can predict the label \hat{y} for some input x , even if it is not in the dataset, i.e. $\hat{y} = f(x)$.

Many possible models

- $f(x) = w_1 x + w_0$.
- $f(x) = w_2 x^2 + w_1 x + w_0$.
- $f(x) = \frac{1}{e^{-(w_1 x + w_0)} + 1}$.
- The numbers w_0 , w_1 and w_2 are called model parameters.
- We often write the model as $f(x; \mathbf{w})$, stacking all parameters to a vector \mathbf{w} .

How would you fit a line?

Can you find a line that passes through $(0, 0)$ and $(1, 1)$?

How would you fit a quadratic curve?

Can you find a quadratic curve that passes through $(0, 0)$, $(1, 1)$ and $(-1, 1)$?

What model do we use for this dataset?

- Open `demo_boston_housing_one_variable.ipynb`
- Can you find a line that go through ALL of the data points?
Why?

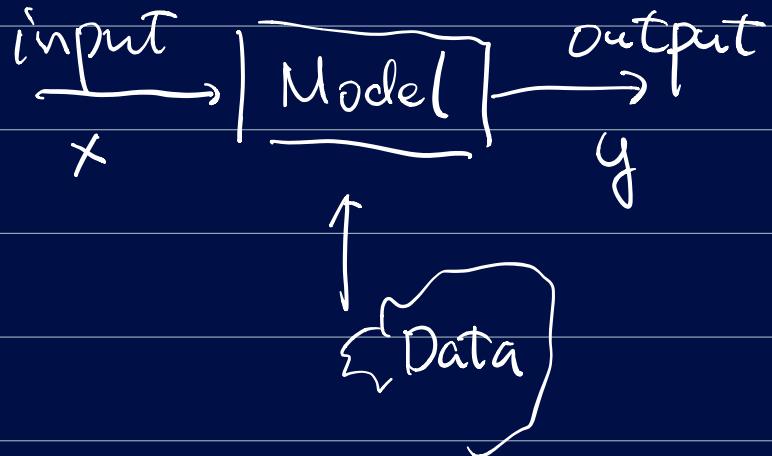
Is Your Model a Good Fit?

- How would you determine if your model is a good fit or not?
 - How will you determine this?
 - Is there a quantitative way?

Error Functions

- An **error function** quantifies the discrepancy between your model and the data.
 - They are non-negative, and go to zero as the model gets better.
- Common Error Functions:
 - Mean Squared Error: $MSE = \frac{1}{N} \sum_{i=1}^N \|y_i - \hat{y}_i\|^2$
 - Mean Absolute Error: $MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$
- In later units, we will refer to these as **cost functions** or **loss functions**.
- Compute MSE on your model
- How do we interpret MSE? MAE?

Machine Learning



Ex: x : movie genre

y : IMDB score

Dataset: $\underbrace{(x_i, y_i)}_{\text{a sample}} \quad i = 1, 2, \dots N$

We also call x features, y labels

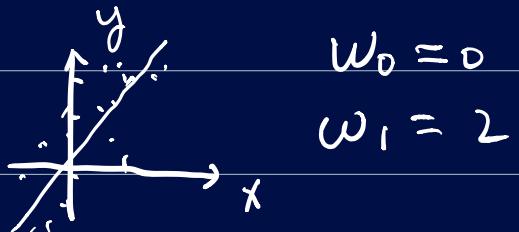
Task: to find a model $f(x)$ to predict

$$\hat{y} = f(x)$$

$(\hat{y}$: "y hat", denotes the prediction given by the model)

$$Ex: \hat{y} = f(x) = \underline{\omega_1}x + \underline{\omega_0}$$

"linear model"



w_0, w_1 : parameters of the model

Cost functions: measure how well the model fits the data

$$MSE: \frac{1}{N} \sum_{i=1}^N \| y_i - \hat{y}_i \|^2$$

y_i : the labels in the dataset
"ground truth" label

$\hat{y}_i = f(x_i)$ predictions

Compute MSE in NumPy

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix} \quad \hat{y} = \begin{bmatrix} \hat{y}_1 \\ \vdots \\ \hat{y}_N \end{bmatrix} \quad \text{MSE: } \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$$y - \hat{y} = \begin{bmatrix} y_1 - \hat{y}_1 \\ y_2 - \hat{y}_2 \\ \vdots \\ y_N - \hat{y}_N \end{bmatrix}$$

In NumPy, $z = (y - \hat{y}) ** 2$, this will square

every element $\begin{bmatrix} (y_1 - \hat{y}_1)^2 \\ \vdots \\ (y_N - \hat{y}_N)^2 \end{bmatrix}$ In our exercise
 $N = 506$

$\text{np.mean}(z)$ will give $\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$

For a vector z with N elements
 $\text{np.mean}(z)$ gives $\frac{1}{N} \sum_{i=1}^N z_i$