

Day 10: Final Project Presentations & Going Further with ML/DL

Summer STEM: Machine Learning

Department of Electrical and Computer Engineering
NYU Tandon School of Engineering
Brooklyn, New York

July 3, 2020

Outline

- 1 Schedule**
- 2 Deep Learning Applications**
- 3 Supervised Learning**
- 4 Unsupervised Learning**
- 5 Social Impact of Machine Learning**
- 6 Course Takeaway**

Schedule

- 1 Finishing Up Projects (10:00-11:30)
- 2 Presentations (11:30-13:00)
- 3 Lunch (13:00-14:30pm)
- 4 Final Lecture: Deep Learning Applications and its Social Impacts (14:30-16:00)
- 5 Goodbye! (16:00 - ∞?)

Outline

- 1 Schedule**
- 2 Deep Learning Applications**
- 3 Supervised Learning**
- 4 Unsupervised Learning**
- 5 Social Impact of Machine Learning**
- 6 Course Takeaway**

Outline

- 1 Schedule**
- 2 Deep Learning Applications**
- 3 Supervised Learning**
- 4 Unsupervised Learning**
- 5 Social Impact of Machine Learning**
- 6 Course Takeaway**


Object Detection

- Faster-RCNN
- YoLo

Object Detection

R-CNN Architecture

R-CNN: Regions with CNN features




Region Proposal Based Object Detection

Object Detection

R-CNN Architecture

R-CNN: Regions with CNN features




Region Proposal Based Object Detection

Object Detection

YOLO

- Divides the image into $n \times n$ grid-cells
- For each grid cell,
 - predicts B bounding boxes and its box confidence score
 - Each box will have its class probability
 - All class probabilities are combined to detect one object

Object Detection



YOLO (<https://pjreddie.com/darknet/yolo/>)

Semantic Segmentation

- Every Pixel is associated with a class
- Encoder-decoder structure
- Decode using transposed convolution or deconvolution




Image Segmentation (Source:
<https://www.jeremyjordan.me/semantic-segmentation/>)

Instance Segmentation





Image generated using Mask-RCNN
(<https://github.com/facebookresearch/detectron2>)

Outline


- 1 Schedule**
- 2 Deep Learning Applications**
- 3 Supervised Learning**
- 4 Unsupervised Learning**
- 5 Social Impact of Machine Learning**
- 6 Course Takeaway**

Autoencoders




From Jeremy Jordan's Post on Autoencoders

Autoencoders



From Jeremy Jordan's Post on Variational Autoencoders

Variational Autoencoders




From Jeremy Jordan's Post on Variational Autoencoders

Variational Autoencoders




VAE face generation implemented by Wojciech Mormul

GANs: Generative Adversarial Networks



Beginning



Generated images (yellow) on a) MNIST b) TFD c) CIFAR-10 (MLP model) d) CIFAR-10 (Conv model)
"Generative Adversarial Networks", Goodfellow et. al. 2014

Progress



Improvement of GANs in producing photo-realistic faces over the years

Celebrity Faces



Human face generation one of the most difficult tasks

Applications of GANs



Image Colorization (Source: <https://github.com/jantic/DeOldify>)

Applications of GANs




Image Synthesis (Source: <https://github.com/NVlabs/SPADEn>)

Applications of GANs




Image Super-Resolution (Source: Dahl et al., "Pixel recursive super resolution")

Outline


- 1 Schedule**
- 2 Deep Learning Applications**
- 3 Supervised Learning**
- 4 Unsupervised Learning**
- 5 Social Impact of Machine Learning**
- 6 Course Takeaway**

How Would You Use ML/DL?

- Think about potential applications with deep learning.
- Discuss its social implications.
- We will dive deeper into this topic later!

Can AI/ML be Biased?

PULSE is a face depixelizing algorithm, but...



Can AI/ML be Biased?

(From the article Design AI so that it's fair)

- When Google Translate converts news articles written in Spanish into English, phrases referring to women often become 'he said' or 'he wrote'.
- Software designed to warn people using Nikon cameras when the person they are photographing seems to be blinking tends to interpret Asians as always blinking.

Other Sources of Bias

Now that we know biased data lead to biased model, are there any other sources of bias in our machine learning pipeline?

De-Bias ML Can be Hard

How would you solve this problem?

The image shows a 2x3 grid of six photographs. The top row contains 'Skyscrapers' (two tall buildings), 'Airplanes' (a view from an airplane window), and 'Cars' (a car parked at night). The bottom row contains 'Bikes' (a person riding a bicycle), 'Gorillas' (two gorillas), and 'Graduation' (a group of people in graduation attire). Each image has a caption below it.

diri noir avec banan @jackyalcine · Jun 29
Google Photos, y'all [REDACTED] My friend's not a gorilla.

813 394 TWITTER

Safety of AI

Boston Dynamics Parkour Atlas: What machine learning algorithms might have been used here?

Safety of AI

- The same model can have drastically different performance for different hyper-parameters.
- 100% accuracy is rarely achieved on unseen data.
- Should we let a medical robot with CNN-based vision system perform surgery autonomously?
- If a self-driving car crashes and hurts people, who should be responsible for it?

Carbon Footprint of Deep Learning

Common carbon footprint benchmarks

in lbs of CO₂ equivalent




Chart: MIT Technology Review • Source: Strubell et al. • Created with Datawrapper

Source: MIT Tech Review

Outline

- 1 Schedule**
- 2 Deep Learning Applications**
- 3 Supervised Learning**
- 4 Unsupervised Learning**
- 5 Social Impact of Machine Learning**
- 6 Course Takeaway**

Course Takeaway

- ML is the combination of math and computer science.
- We've only shown you a subsection
 - Supervised Learning: Linear/Logistic Regression and Neural Networks
- Deep learning has wide applications, but we are also responsible for its consequences. —The greater the power, the greater the responsibility!

Thank You!

■ Thank You!